

Andreas Holzinger
185.A83 Machine Learning for Health Informatics
2019S, VU, 2.0 h, 3.0 ECTS
Lecture 04 – Dienstag, 02.04.2019



From Decision Making under uncertainty to graphical models and MCMC

andreas.holzinger AT tuwien.ac.at

<https://hci-kdd.org/machine-learning-for-health-informatics-class-2019>



- 00 Reflection from last lecture
- 01 Decision Making under uncertainty
- 02 Some Basics of Markov Processes
- 03 Some Basics of Concept Learning
- 04 Some Basics of Graphs/Networks and Challenges
- 05 Bayes Nets
- 06 Probabilistic Programming
- 07 Markov Chain Monte Carlo (MCMC)
- 08 Metropolis Hastings Algorithm

01 Reflection



Why Explainability? Why Causability?


Causability :=
a property of a person
(Human Intelligence)

Explainability :=
a property of a system
(Artificial Intelligence)

<https://onlinelibrary.wiley.com/doi/full/10.1002/widm.1312>



Andreas Holzinger et al. 2019. Causability and Explainability of AI in Medicine. Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery, doi:10.1002/widm.1312.

*Why did the algorithm do that?
Can I trust these results?
How can I correct an error?*



Input data

A possible solution



Input data


The domain expert can understand why ...

The domain expert can learn and correct errors ...

The domain expert can re-enact on demand ...

- 1) Gradients
- 2) Sensitivity Analysis
- 3) Decomposition Relevance Propagation
 - Pixel-RP, Layer-RP, Deep Taylor Decomposition, ...
- 4) Optimization
 - Local Interpretable Model-Agnostic Explanations (LIME)
 - Black Box Explanations tr. Transparent Approximations (BETA)
- 5) Deconvolution and Guided Backpropagation
- 6) Concept Activation Vectors CAV

Andreas Holzinger LV 706.315 From explainable AI to Causability, 3 ECTS course at Graz University of Technology
<https://hci-kdd.org/explainable-ai-causability-2019> (course given since 2016)



Physician treating a patient
approx. 480 B.C.
Beazley (1963), Attic Red-figured
Vase-Painters, 813, 96.
Department of Greek, Etruscan
and Roman Antiquities, Sully, 1st
floor, Campana Gallery, room 43
Louvre, Paris

Elwyn, G., Edwards, A., Eccles, M. & Rovner, D. 2001. Decision analysis in patient care. The Lancet, 358, (9281), 571-574.



01 Decision Making under uncertainty

Laplace, P.-S. 1781. Mémoire sur les probabilités. *Mémoires de l'Académie Royale des sciences de Paris*, 1778, 227-332.



... permanent decision making
under uncertainty!

d ... data

$\mathcal{H} \dots \{H_1, H_2, \dots, H_n\}$

$\forall h, d \dots$

h ... hypotheses


$$p(h|d) = \frac{p(d|h) * p(h)}{\sum_{h \in H} p(d|h') p(h')}$$


Posterior Probability

Likelihood

Prior Probability



Problem in $\mathbb{R}^n \rightarrow$ complex






Wickens, C. D. (1984) *Engineering psychology and human performance*.
Columbus (OH), Charles Merrill, modified by Holzinger, A.

You are talking to your colleague and want to refer to the middle object – which wording would you prefer: circle or blue?




Frank, M. C. & Goodman, N. D. 2012. Predicting pragmatic reasoning in language games. *Science*, 336, (6084), 998-998, doi:10.1126/science.1218633.

```
var literalListener = function(property){  
  Infer(function(){  
    var object = refPrior(context)  
    condition(object[property])  
    return object  
  })  
  
  var speaker = function(object) {  
    Infer(function(){  
      var property = propPrior()  
      condition(  
        object ==  
        sample(speaker(object)))  
      return object  
    })  
  
    var listener = function(property) {  
      Infer(function(){  
        var object = refPrior(context)  
        condition(utterance ==  
          sample(speaker(object)))  
        return object  
      })  
    }  
  }  
}
```





Goodman, N. D. & Frank, M. C. 2016. Pragmatic language interpretation as probabilistic inference. Trends in Cognitive Sciences, 20, (11), 818-829.

- PGM can be seen as a combination between
 - **Graph Theory + Probability Theory + Machine Learning**
 - One of the most exciting advancements in AI in the last decades – with enormous future potential
 - Compact representation for exponentially-large probability distributions
 - Example Question:
“Is there a path connecting two proteins?”
-
- $\text{Path}(X, Y) := \text{edge}(X, Y)$
 - $\text{Path}(X, Y) := \text{edge}(X, Y), \text{path}(Z, Y)$
 - This can NOT be expressed in first-order logic
 - Would need a Turing-complete fully-fledged language



2) Some basics of Markov Processes in Machine Learning


- Markov processes are ...
- random processes in which the future, given the present, is independent of the past!
- one of the most important classes of random processes!



Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., Graves, A., Riedmiller, M., Fidjeland, A. K., Ostrovski, G., Petersen, S., Beattie, C., Sadik, A., Antonoglou, I., King, H., Kumaran, D., Wierstra, D., Legg, S. & Hassabis, D. 2015. Human-level control through deep reinforcement learning. *Nature*, 518, (7540), 529-533, doi:10.1038/nature14236

YouTube DE

reinforcement learning space invaders



Deep Q network playing Space Invaders



eldubro

Subscribe 11

1,855

Add to

Share

More

1 0

Up next







Image credit to David Silver, UCL




```
initialize  $V(s)$  arbitrarily
loop until policy good enough
    loop for  $s \in \mathcal{S}$ 
        loop for  $a \in \mathcal{A}$ 
             $Q(s, a) := R(s, a) + \gamma \sum_{s' \in \mathcal{S}} T(s, a, s')V(s')$ 
             $V(s) := \max_a Q(s, a)$ 
        end loop
    end loop
```

Kaelbling, L. P., Littman, M. L. & Moore, A. W. 1996. Reinforcement learning: A survey.
Journal of Artificial Intelligence Research, 4, 237-285.


```
for t = 1, ..., n do
    The agent perceives state  $s_t$ 
    The agent performs action  $a_t$ 
    The environment evolves to  $s_{t+1}$ 
    The agent receives reward  $r_t$ 
end for
```

Intelligent behavior arises from the actions of an individual seeking to **maximize its received reward signals** in a **complex and changing world**




Sutton, R. S. & Barto, A. G. 1998. Reinforcement learning: An introduction, Cambridge MIT press

- Supervised:
Learner told best a
- Exhaustive:
Learner shown every possible x
- One-shot: Current x independent of past a



Littman, M. L. 2015. Reinforcement learning improves behaviour from evaluative feedback. Nature, 521, (7553), 445-451.

- Markov decision processes specify setting and tasks
- Planning methods use knowledge of P and R to compute a good policy π
- Markov decision process model captures both sequential feedback and the more specific one-shot feedback (when $P(s'|s, a)$ is independent of both s and a)



$$Q^*(s, a) = R(s, a) + \gamma \sum P(s'|s, a) \max_{a'} Q^*(s', a')$$

Littman, M. L. 2015. Reinforcement learning improves behaviour from evaluative feedback. Nature, 521, (7553), 445-451.

- 1) Observes
- 2) Executes
- 3) Receives Reward
- Executes action A_t :
- $O_t = s a_t = s e_t$
- Agent state =
environment state =
information state
- Markov decision
process (MDP)






Image credit to David Silver, UCL

- i.e. whatever data the environment uses to pick the next observation/reward
- The environment state is not usually visible to the agent
- Even if S is visible, it may contain irrelevant information
- A State S_t is Markov iff:



$$\mathbb{P}[S_{t+1}|S_t] = \mathbb{P}[S_{t+1}|S_1, \dots, S_t]$$

- i.e. whatever information the agent uses to pick the next action
- it is the information used by reinforcement learning algorithms
- It can be any function of history:
- $S = f(H)$



$$H_t = O_1, R_1, A_1, \dots, A_{t-1}, O_t, R_t$$

- RL agent components:
 - Policy: agent's behaviour function
 - Value function: how good is each state and/or action
 - Model: agent's representation of the environment
- Policy as the agent's behaviour
 - is a map from state to action, e.g.
 - Deterministic policy: $a = \pi(s)$
 - Stochastic policy: $\pi(a|s) = P[A_t = a | S_t = s]$
- Value function is prediction of future reward:

$$v_\pi(s) = \mathbb{E}_\pi [R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s]$$


- Partial observability: when agent only indirectly observes environment (robot which is not aware of its current location; good example: Poker play: only public cards are observable for the agent):
- Formally this is a Partially Observable Markov Decision Process (POMDP):
 - Agent must construct its own state representation S , for example:
 - Complete history: $S_t^a = H_t$
 - Beliefs of environment state: $S_t^a = (\mathbb{P}[S_t^e = s^1], \dots, \mathbb{P}[S_t^e = s^n])$
 - Recurrent neural network: $S_t^a = \sigma(S_{t-1}^a W_s + O_t W_o)$

3) Some basics of Concept Learning

- **Deductive Reasoning** = Hypothesis > Observations > Logical Conclusions (general → specific – proven correctness)
 - DANGER: Hypothesis must be correct! DR defines whether the truth of a conclusion can be determined for that rule, based on the truth of premises: A=B, B=C, therefore A=C
- **Inductive reasoning** = makes broad generalizations from specific observations (specific → general – not proven correctness)
 - DANGER: allows a conclusion to be false if the premises are true
 - generate hypotheses and use DR for answering specific questions
- **Abductive reasoning** = inference = to get the best explanation from an incomplete set of preconditions.
 - Given a true conclusion and a rule, it attempts to select some possible premises that, if true also, may support the conclusion, though not uniquely.
 - Example: "When it rains, the grass gets wet. The grass is wet. Therefore, it might have rained." This kind of reasoning can be used to develop a hypothesis, which in turn can be tested by additional reasoning or data.

- Bruner, Goodnow, and Austin (1956) published “A Study of Thinking”, which became a landmark in cognitive science and has much influence on machine learning.

- Rule-Based Categories
- A concept specifies conditions for membership



Jerome S. Bruner, Jacqueline J. Goodnow & George A. Austin 1986. A Study of Thinking, Transaction Books.



Salakhutdinov, R., Tenenbaum, J. & Torralba, A. 2012. One-shot learning with a hierarchical nonparametric Bayesian model. *Journal of Machine Learning Research*, 27, 195-207.


Quaxl

Quaxl


Quaxl



Salakhutdinov, R., Tenenbaum, J. & Torralba, A. 2012. One-shot learning with a hierarchical nonparametric Bayesian model. *Journal of Machine Learning Research*, 27, 195-207.





$$P(h|d) = \frac{P(d|h)P(h)}{\sum_{h'} P(d|h')P(h')} \propto P(d|h)P(h)$$



Tenenbaum, J. B., Kemp, C., Griffiths, T. L. & Goodman, N. D. 2011. How to grow a mind: Statistics, structure, and abstraction. *Science*, 331, (6022), 1279-1285.


- which is highly relevant for ML research, concerns the factors that determine the subjective difficulty of concepts:
- Why are some concepts psychologically extremely simple and easy to learn,
- while others seem to be extremely difficult, complex, or even incoherent?
- These questions have been studied since the 1960s but are still unanswered ...

Feldman, J. 2000. Minimization of Boolean complexity in human concept learning. *Nature*, 407, (6804), 630-633, doi:10.1038/35036586.




- Cognition as probabilistic inference
 - Visual perception, language acquisition, motor learning, associative learning, memory, attention, categorization, reasoning, causal inference, decision making, theory of mind
- Learning concepts from examples
- Learning causation from correlation
- Learning and applying intuitive theories (balancing complexity vs. fit)

- Similarity
- Representativeness and evidential support
- Causal judgement
- Coincidences and causal discovery
- Diagnostic inference
- Predicting the future




Tenenbaum, J. B., Griffiths, T. L. & Kemp, C. 2006. Theory-based Bayesian models of inductive learning and reasoning. Trends in cognitive sciences, 10, (7), 309-318.



Lake, B. M., Salakhutdinov, R. & Tenenbaum, J. B. 2015. Human-level concept learning through probabilistic program induction. *Science*, 350, (6266), 1332-1338, doi:10.1126/science.aab3050.

A Bayesian program learning (BPL) framework, capable of learning a large class of visual concepts from just a single example and generalizing in ways that are mostly indistinguishable from people



Lake, B. M., Salakhutdinov, R. & Tenenbaum, J. B. 2015. Human-level concept learning through probabilistic program induction. *Science*, 350, (6266), 1332-1338,

4) Graphs=Networks





Image from <https://people.kth.se/~carlofi/teaching/FEL3250-2013/courseinfo.html>



Pearl, J. 1988. Embracing causality in default reasoning. Artificial Intelligence, 35, (2), 259-271.

A.M. TURING CENTENARY CELEBRATION WEBCAST

The screenshot shows the homepage of the ACM A.M. Turing Award website. At the top left is the ACM logo. To its right is a large blue banner with the text "A.M. TURING AWARD". Below the banner is a grid of 24 small portraits of Turing award winners. Above the grid is a search bar with the placeholder "TYPE HERE". Below the grid are three navigation links: "ALPHABETICAL LISTING", "YEAR OF THE AWARD", and "RESEARCH SUBJECT".

JUDEA PEARL

United States – 2011

CITATION

For fundamental contributions to artificial intelligence through the development of a calculus for probabilistic and causal reasoning.

SHORT ANNOTATED BIBLIOGRAPHY ACM DL AUTHOR PROFILE ACM TURING AWARD LECTURE VIDEO RESEARCH SUBJECTS ADDITIONAL MATERIALS

Judea Pearl created the representational and computational foundation for the processing of information under uncertainty.

He is credited with the invention of *Bayesian networks*, a mathematical formalism for defining complex probability models, as well as the principal algorithms used for inference in these models. This work not only revolutionized the field of artificial intelligence but also became an important tool for many other branches of engineering and the natural sciences. He later created a mathematical framework for *causal inference* that has had significant impact in the social sciences.

Judea Pearl was born on September 4, 1936, in Tel Aviv, which was at that time administered under the British Mandate for Palestine. He grew up in *Bnei Brak*, a Biblical town his grandfather went to reestablish in 1924. In 1956, after serving in the Israeli army and joining a Kibbutz, Judea decided to study engineering. He attended the Technion, where he met his wife, Ruth, and received a B.S. degree in Electrical Engineering in 1960. Recalling the Technion faculty members in a 2012 interview in the *Technion Magazine*, he emphasized the thrill of learning.

http://amturing.acm.org/vp/pearl_2658896.cfm



Scientific Background on the Nobel Prize in Chemistry 2013

DEVELOPMENT OF MULTISCALE MODELS FOR
COMPLEX CHEMICAL SYSTEMS



Photo: A. Mahmoud

Martin Karplus

Prize share: 1/3



Photo: A. Mahmoud

Michael Levitt

Prize share: 1/3




Photo: A. Mahmoud

Arieh Warshel

Prize share: 1/3

http://www.nobelprize.org/nobel_prizes/chemistry/laureates/2013



http://news.harvard.edu/gazette/story/2013/10/nobel_prize_awarded_2013/

- **Graphs as models for networks**
- given as direct input (point cloud data sets)
- Given as properties of a structure
- Given as a representation of information (e.g. Facebook data, viral marketing, etc., ...)
- **Graphs as nonparametric basis**
- we learn the structure from samples and infer
- flat vector data, e.g. similarity graphs
- encoding structural properties (e.g. smoothness, independence, ...)

We skip this interesting chapter for now ...



Network challenges

NGC 5139 Omega Centauri by Edmund Halley in 1677, ESO, Atacama, Chile

Time



e.g. Entropy




Dali, S. (1931) The persistence of memory

Space


e.g. Topology



Bagula & Bourke (2012) Klein-Bottle




P versus NP and the Computational Complexity Zoo, please have a look at
<https://www.youtube.com/watch?v=YX40hbAHx3s>




ATOM	1	N	GLY	A	1	44.842	51.034	101.284	0.01	27.20
ATOM	2	CA	GLY	A	1	45.640	50.230	100.389	0.01	26.99
ATOM	3	C	GLY	A	1	46.692	49.648	101.308	0.01	26.80
ATOM	4	O	GLY	A	1	46.895	50.222	102.381	0.01	26.91
ATOM	5	N	SER	A	2	47.283	48.516	100.951	1.00	26.26
ATOM	6	CA	SER	A	2	48.277	47.866	101.761	1.00	26.17
ATOM	7	C	SER	A	2	49.212	47.031	100.845	1.00	24.21
ATOM	8	O	SER	A	2	49.060	47.195	99.630	1.00	19.77
ATOM	9	CB	SER	A	2	47.438	47.091	102.800	1.00	26.31
ATOM	10	OG	SER	A	2	46.276	46.356	102.404	1.00	27.99
ATOM	11	N	HIS	A	3	50.147	46.186	101.370	1.00	23.93
ATOM	12	CA	HIS	A	3	51.129	45.389	100.609	1.00	21.44
ATOM	13	C	HIS	A	3	50.953	43.905	100.849	1.00	20.32
ATOM	14	O	HIS	A	3	50.530	43.595	101.950	1.00	22.00
ATOM	15	CB	HIS	A	3	52.555	45.674	100.990	1.00	19.69
ATOM	16	CG	HIS	A	3	52.940	47.090	100.611	1.00	21.44
ATOM	17	ND1	HIS	A	3	53.371	47.470	99.422	1.00	20.87
ATOM	18	CD2	HIS	A	3	52.956	48.175	101.433	1.00	21.69
ATOM	19	CE1	HIS	A	3	53.676	48.730	99.476	1.00	20.57

Wiltgen, M. & Holzinger, A. (2005) Visualization in Bioinformatics: Protein Structures with Physicochemical and Biological Annotations. In: *Central European Multimedia and Virtual Reality Conference. Prague, Czech Technical University (CTU)*, 69-74



Wiltgen, M., Holzinger, A. & Tilz, G. P. (2007) Interactive Analysis and Visualization of Macromolecular Interfaces Between Proteins. In: *Lecture Notes in Computer Science (LNCS 4799)*. Berlin, Heidelberg, New York, Springer, 199-212.



Nodes = proteins
Links = physical interactions
(bindings)
Red Nodes = lethal
Green Nodes = non-lethal
Orange = slow growth
Yellow = not known


Jeong, H., Mason, S. P., Barabasi, A. L. & Oltvai, Z. N. (2001) Lethality and centrality in protein networks. *Nature*, 411, 6833, 41-42.


Light blue = known proteins

Orange = disease proteins


Yellow ones = not known yet

Stelzl, U. et al.
(2005) A Human
Protein-Protein
Interaction
Network: A
Resource for
Annotating the
Proteome. *Cell*,
122, 6, 957-968.






Hurst, M. (2007), Data Mining: Text Mining, Visualization and Social Media. Online available: http://datamining.typepad.com/data_mining/2007/01/the_blogosphere.html, last access: 2011-09-24



Aral, S. (2011)
Identifying Social
Influence: A Comment
on Opinion Leadership
and Social Contagion in
New Product Diffusion.
Marketing Science, 30,
2, 217-223.



Barabási, A. L.,
Gulbahce, N. &
Loscalzo, J. 2011.
Network medicine: a
network-based
approach to human
disease. *Nature Reviews
Genetics*, 12, 56-68.

05 Bayesian

Networks

“Bayes’ Nets”

- is a **probabilistic model**, consisting of two parts:
- 1) a dependency structure and
- 2) local probability models.

$$p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i \mid Pa(x_i))$$

Where $Pa(x_i)$ are the parents of x_i


BN inherently model the uncertainty in the data. They are a successful marriage between probability theory and graph theory; allow to model a multidimensional probability distribution in a sparse way by searching independency relations in the data. Furthermore this model allows different strategies to integrate two data sources.


Pearl, J. (1988) *Probabilistic reasoning in intelligent systems: networks of plausible inference*. San Francisco, Morgan Kaufmann.

$$p(X_1, \dots, X_7) =$$

$$p(X_1)p(X_2)p(X_3)p(X_4|X_1, X_2, X_3) \cdot$$


$$p(X_5|X_1, X_3)p(X_6|X_4)p(X_7|X_4, X_5)$$






Overmoyer, B. A.,
Lee, J. M. &
Lerwill, M. F.
(2011) Case 17-
2011 A 49-Year-
Old Woman with a
Mass in the Breast
and Overlying Skin
Changes. *New
England Journal of
Medicine*, 364, 23,
2246-2254.

- = the prediction of the future course of a disease conditional on the patient's history and a projected treatment strategy
- Danger: probable Information !
- Therefore valid prognostic models can be of great benefit for clinical decision making and of great value to the patient, e.g., for notification and quality of-life decisions




Knaus, W. A., Wagner, D. P. & Lynn, J. (1991) Short-term mortality predictions for critically ill hospitalized adults: science and ethics. *Science*, 254, 5030, 389.



van Gerven, M. A. J., Taal, B. G. & Lucas, P. J. F. (2008) Dynamic Bayesian networks as prognostic models for clinical patient management. *Journal of Biomedical Informatics*, 41, 4, 515-529.


Category	Node description	State description
Diagnosis	Breast cancer	Present, absent.
Clinical history	Habit of drinking alcoholic beverages and smoking	Yes, no.
	Taking female hormones	Yes, no.
	Have gone through menopause	Yes, no.
	Have ever been pregnant	Yes, no.
	Family member has breast cancer	Yes, no.
Physical findings	Nipple discharge	Yes, no.
	Skin thickening	Yes, no.
	Breast pain	Yes, no.
	Have a lump(s)	Yes, no.
Mammographic findings	Architectural distortion	Present, absent.
	Mass	Score from one to three, score from four to five, absent
	Microcalcification cluster	Score from one to three, score from four to five, absent
	Asymmetry	Present, absent.

Wang, X. H., et al. (1999) Computer-assisted diagnosis of breast cancer using a data-driven Bayesian belief network. *International Journal of Medical Informatics*, 54, 2, 115-126.




Wang, X. H., et al. (1999) Computer-assisted diagnosis of breast cancer using a data-driven Bayesian belief network. *International Journal of Medical Informatics*, 54, 2, 115-126.


- Integrating microarray data from multiple studies to increase sample size;
- = approach to the development of more robust prognostic tests



Xu, L., Tan, A., Winslow, R. & Geman, D. (2008) Merging microarray data from separate breast cancer studies provides a robust prognostic test. *BMC Bioinformatics*, 9, 1, 125-139.



Gevaert, O., Smet, F. D., Timmerman, D., Moreau, Y. & Moor, B. D. (2006) Predicting the prognosis of breast cancer by integrating clinical and microarray data with Bayesian networks. *Bioinformatics*, 22, 14, 184-190.



Gevaert, O., Smet, F. D.,
Timmerman, D.,
Moreau, Y. & Moor, B. D.
(2006) Predicting the
prognosis of breast
cancer by integrating
clinical and microarray
data with Bayesian
networks.
Bioinformatics, 22, 14,
184-190.

- First the structure is learned using a search strategy.
- Since the number of possible structures increases super exponentially with the number of variables,
- the well-known greedy search algorithm K2 can be used in combination with the Bayesian Dirichlet (BD) scoring metric:

$$p(S|D) \propto p(S) \prod_{i=1}^n \prod_{j=1}^{q_i} \left[\frac{\Gamma(N'_{ij})}{\Gamma(N'_{ij} + N_{ij})} \prod_{k=1}^{r_i} \frac{\Gamma(N'_{ijk} + N_{ijk})}{\Gamma(N'_{ijk})} \right]$$

N_{ijk} ... number of cases in the data set D
having variable i in state k associated with the j -th instantiation
of its parents in current structure S .
 n is the total number of variables.

- Next, N_{ij} is calculated by summing over all states of a variable:
- $N_{ij} = \sum_{k=1}^{r_i} N_{ijk} \cdot N'_{ijk}$ and N'_{ij} have similar meanings but refer to prior knowledge for the parameters.
- When no knowledge is available they are estimated using $N_{ijk} = N/(r_i q_i)$
- with N the equivalent sample size,
- r_i the number of states of variable i and
- q_i the number of instantiations of the parents of variable i .
- $\Gamma(\cdot)$ corresponds to the gamma distribution.
- Finally $p(S)$ is the prior probability of the structure.
- $p(S)$ is calculated by:
- $p(S) = \prod_{i=1}^n \prod_{l_i=1}^{p_i} p(l_i \rightarrow x_i) \prod_{m_i=1}^{o_i} p(m_i x_i)$
- with p_i the number of parents of variable x_i and o_i all the variables that are not a parent of x_i .
- Next, $p(a \rightarrow b)$ is the probability that there is an edge from a to b while $p(ab)$ is the inverse, i.e. the probability that there is no edge from a to b


- Estimating the parameters of the local probability models corresponding with the dependency structure.
- CPTs are used to model these local probability models.
- For each variable and instantiation of its parents there exists a CPT that consists of a set of parameters.
- Each set of parameters was given a uniform Dirichlet prior:

$$p(\theta_{ij}|S) = \text{Dir}(\theta_{ij}|N'_{ij1}, \dots, N'_{ijk}, \dots, N'_{ijr_i})$$

Note: With θ_{ij} a parameter set where i refers to the variable and j to the j -th instantiation of the parents in the current structure. θ_{ij} contains a probability for every value of the variable x_i given the current instantiation of the parents. Dir corresponds to the Dirichlet distribution with $(N'_{ij1}, \dots, N'_{ijr_i})$ as parameters of this Dirichlet distribution. Parameter learning then consists of updating these Dirichlet priors with data. This is straightforward because the multinomial distribution that is used to model the data, and the Dirichlet distribution that models the prior, are conjugate distributions. This results in a Dirichlet posterior over the parameter set:

$$p(\theta_{ij}|D, S) = \text{Dir}(\theta_{ij}|N'_{ij1} + N_{ij1}, \dots, N'_{ijk} + N_{ijk}, \dots, N'_{ijr_i} + N_{ijr_i})$$

with N_{ijk} defined as before.



Gevaert, O., Smet, F. D.,
Timmerman, D., Moreau, Y. &
Moor, B. D. (2006) Predicting
the prognosis of breast cancer
by integrating clinical and
microarray data with Bayesian
networks. *Bioinformatics*, 22,
14, 184-190.


- For certain cases it is tractable if:
 - Just one variable is unobserved
 - We have singly connected graphs (no undirected loops -> belief propagation)
 - Assigning probability to fully observed set of variables
- Possibility: Monte Carlo Methods (generate many samples according to the Bayes Net distribution and then count the results)
- Otherwise: approximate solutions, NOTE:
Sometimes it is better to have an approximate solution to a complex problem – than a perfect solution to a simplified problem

Often it is better to
have a good solution
within time – than an
perfect solution
(much) later ...



06 Probabilistic Programming

- C → Probabilistic-C
- Scala → Figaro
- Scheme → Church
- Excel → Tabular
- Prolog → Problog
- Javascript → webPP
- → Venture
- Python → PyMC



PyMC Pythonic Markov chain Monte Carlo

Probabilistic Program	Graphical Model
Variables	Variable nodes
Functions/operators	Factor nodes/edges
Fixed size loops/arrays	Plates
If statements	Gates (Minka & Winn)
Variable sized loops, Complex indexing, jagged arrays, mutation, recursion, objects/ properties...	No common equivalent

Sequence	Outcome
CGTCGGAGGTACATGATTGAAAGAAAAACCT	Y
GCGCCTTGACATCTCTTAAATCTCAGTCA	X
TTAAAATAGCAGAGACACTTCTACTGATACT	Y
CCAAGAGCCTCGTAATTAAGTATTGCAATA	Y
TTATGACGTCGTTCGAGTGGATTTGTCTT	X
...	...

1

- Simple example: Nucleotide "A" may follow nucleotide "T" in the sequences more frequently for outcome X than for outcome Y,

$$P(A|T, X) > P(A|T, Y) \quad 2$$



- Compute maximum a posteriori estimates of the probabilities:

```
from pymc import MAP, Model
model = Model({'f_X': f_X, 'prob_dist': prob_dist})
M = MAP(model)
M.fit() # Helder-Mead Optimization
```

- The MAP estimates are now contained in the M.prob_dist value:

```
>>> print M.prob_dist.value
[ 0.1947259, 0.26842748, 0.23265728]
```

$$P(\theta|D) = \frac{P(D|\theta) \cdot P(\theta)}{P(D)}$$

6


- Specify the value to maximize using numerical simulation, as well as the expected form of the posterior distribution:

```
from pymc import Categorical
f_X = Categorical('cat', prob_dist, value=exp_data, observed=True)
```

$$P(\theta|D) = \frac{P(D|\theta) \cdot P(\theta)}{P(D)}$$

- Specify the prior distribution:

```
import numpy as np
from pymc import Dirichlet # conjugate prior
alpha = np.array([30.0, 25.0, 20.0, 25.0])
prob_dist = Dirichlet('prob_dist', alpha)
```



3

$$P(\theta|D) = \frac{P(D|\theta) \cdot P(\theta)}{P(D)}$$

- Specify the experimental data:

```
exp_data = np.array([1, 1, 3, 2, 2, 1, 0, ...])
```

Experimental Data

Observation #	Nucleotide
1	1
2	1
3	3
4	2
5	2
6	1
7	0

4

$$P(\theta|D) = \frac{P(D|\theta) \cdot P(\theta)}{P(D)}$$

Image Source: Dan Williams, Life Technologies, Austin TX

07 Markov Chain Monte Carlo (MCMC)






Monte Carlo Method (MC)

Monte Carlo Sampling

Markov Chains (MC)

MCMC

Metropolis-Hastings



- often we want to calculate characteristics of a **high-dimensional** probability distribution ... $p(\mathcal{D}|\theta)$

$$p(h|d) \propto p(\mathcal{D}|\theta) * p(h)$$


Posterior integration problem: (almost) all statistical inference can be deduced from the posterior distribution by calculating the appropriate sums, which involves an integration:

$$J = \int f(\theta) * p(\theta|\mathcal{D}) d\theta$$

- **Statistical physics:** computing the partition function – this is evaluating the posterior probability of a hypothesis and this requires summing over all hypotheses ... remember:

$$\mathcal{H} = \{H_1, H_2, \dots, H_n\} \quad \forall(h, d)$$

$$P(h|d) = \frac{P(d|h) * P(h)}{\sum_{h' \in \mathcal{H}} P(d|h')P(h')}$$








- Class of algorithms that rely on **repeated random sampling**
- Basic idea: using **randomness** to solve problems with high uncertainty (Laplace, 1781)
- For solving **multidimensional integrals** which would otherwise intractable
- For simulation of systems with **many dof**
- e.g. fluids, gases, particle collectives, **cellular structures** - see our last tutorial on Tumor growth simulation!



- for solving problems of probabilistic inference involved in developing computational models
- as a source of hypotheses about how the human mind might solve problems of inference
- For a function $f(x)$ and distribution $P(x)$, the expectation of f with respect to P is generally the average of f , when x is drawn from the probability distribution $P(x)$

$$\mathbb{E}_{p(x)}(f(x)) = \sum_X f(x)P(x)dx$$

- Solving intractable integrals
- Bayesian statistics: **normalizing** constants, expectations, marginalization
- Stochastic Optimization
- Generalization of simulated annealing
- Monte Carlo expectation maximization (EM)



- Physical simulation
- estimating neutron diffusion time
- Computing expected utilities and best responses toward Nash equilibria
- Computing volumes in high-dimensions
- Computing eigen-functions and values of operators (e.g. Schrödinger)
- Statistical physics
- Counting many things as fast as possible



- Expectation of a function $f(x, y)$ with respect to a random variable x is denoted by $\mathbb{E}_x [f(x, y)]$
- In situations where there is no ambiguity as to which variable is being averaged over, this will be simplified by omitting the suffix, for instance $\mathbb{E}x$.
- If the distribution of x is conditioned on another variable z , then the corresponding conditional expectation will be written $\mathbb{E}_x[f(x)|z]$
- Similarly, the variance is denoted $\text{var}[f(x)]$, and for vector variables the covariance is written $\text{cov}[x, y]$

$$\underset{x}{\operatorname{argmax}} f(x)$$

Normalization: $p(x|y) = \frac{p(y|x) * p(x)}{\int_X p(y|x) * p(x) dx}$

Marginalization: $p(x) = \int_Z p(x, z) dz$

Expectation: $\mathbb{E}_{p(x)}(f(x)) = \int_X f(x)p(x)dx$

08 Metropolis- Hastings Algorithm







Image Source: Peter Mueller,
Anderson Cancer Center

Barber, D. 2012. Bayesian reasoning and machine learning, Cambridge, Cambridge University Press, p. 500

```
1: Choose a starting point  $x^1$ .  
2: for  $i = 2$  to  $L$  do  
3:   Draw a candidate sample  $x^{cand}$  from the proposal  $\tilde{q}(x'|x^{l-1})$ .  
4:   Let  $a = \frac{\tilde{q}(x^{l-1}|x^{cand})p(x^{cand})}{\tilde{q}(x^{cand}|x^{l-1})p(x^{l-1})}$   
5:   if  $a \geq 1$  then  $x^l = x^{cand}$   
6:   else  
7:     draw a random value  $u$  uniformly from the unit interval  $[0, 1]$ .  
8:     if  $u < a$  then  $x^l = x^{cand}$   
9:     else  
10:       $x^l = x^{l-1}$   
11:    end if  
12:  end if  
13: end for
```



- Importance sampling is a technique to approximate averages with respect to an intractable distribution $p(x)$.
- The term ‘sampling’ is arguably a misnomer since the method does not attempt to draw samples from $p(x)$.
- Rather the method draws samples from a simpler importance distribution $q(x)$ and then reweights them
- such that averages with respect to $p(x)$ can be approximated using the samples from $q(x)$.

- The Gibbs Sampler is an interesting special case of MH:







Image Source: Peter Mueller,
Anderson Cancer Center





Azizi, E., Airola, E. M. & Galagan, J. E. 2014.
Learning Modular Structures from Network Data
and Node Variables. Proceedings of the 31st
International Conference on Machine Learning
(ICML). Beijing: JMLR. 1440-1448.




Algorithm 1 RJMCMC for sampling parameters**Inputs:**Node Variables Data \mathbf{X} Network Data \mathbf{B} **for** iterations $j = 1$ to J **do** Sample $\mathcal{A}^{(j+1)}$ given $\mathcal{A}^{(j)}$ using Alg 2 in (Azizi et al., 2014) Sample $\mathcal{S}^{(j+1)}$ given $\mathcal{S}^{(j)}$ using Alg 3 in (Azizi et al., 2014) **for** modules $k = 1$ to $K^{(j)}$ **do** Propose $w_k^{(j+1)} \sim \mathcal{N}(w_k^{(j)}, I)$ Accept with probability P_{mh} ; update $\Sigma^{(j+1)}$ **for** parents $r = 1$ to R_k **do** Propose $z_k^{r(j+1)} \sim \mathcal{N}(z_k^{r(j)}, I)$; accept with P_{mh} Propose $\pi_k^{r(j+1)} \sim \mathcal{N}(\pi_k^{r(j)}, I)$; accept with P_{mh} **end for** **end for** **for** condition $c = 1$ to C **do** Propose $\mu_c^{\mathbf{R}(j+1)} \sim \mathcal{N}(\mu_c^{\mathbf{R}(j)}, I)$; accept with P_{mh} Propose $\gamma_c^{\mathbf{R}(j+1)} \sim \mathcal{N}(\gamma_c^{\mathbf{R}(j)}, I)$; accept with P_{mh} **end for****end for**

Azizi, E., Airolldi, E. M. & Galagan, J. E. 2014. Learning Modular Structures from Network Data and Node Variables.

Proceedings of the 31st International Conference on Machine Learning (ICML). Beijing: JMLR. 1440-1448.
Holzinger Group hci-kdd.org



Azizi, E., Airoldi, E. M. & Galagan, J. E. 2014. Learning Modular Structures from Network Data and Node Variables. Proceedings of the 31st International Conference on Machine Learning (ICML). Beijing: JMLR. 1440-1448.



Henao, R., Lu, J. T., Lucas, J. E., Ferranti, J. & Carin, L. 2016. Electronic health record analysis via deep poisson factor models. Journal of Machine Learning Research JMLR, 17, 1-32.




Thank you!

Questions


- What is the main difference between the ideas of Pierre Simon de Laplace and Lady Lovelace?
- What is medical action consisting most of the time?
- How does a human make a decision - as far as we know?
- What is the main idea of a probabilistic programming language?
- Why did Judea Pearl receive the Turing Award (Nobel Prize in Computer Science)?
- What fields are coming together in PGM?
- What are the challenges in network structures?
- Give a classification of Graphical Models!
- What are plates and nested plates?
- Provide corresponding examples of metabolic networks!

- What is a factored graph?
- Describe the protein structure prediction problem! Why is it hard?
- Why are protein-protein interactions so important?
- Describe the problem of graph-isomorphism!
- How does a Bayes Net work?
- Why is predicting important in clinical medicine?
- What is a Markov-Blankett?
- Which two tasks do we have in Graphical Model Learning?
- Why would we need probabilistic programming languages?
- Describe the main idea of MCMC!
- What is the main problem in marginalization?
- What is the benefit of the MH Algorithm?


Appendix




Bishop, C. M. 2007. Pattern Recognition and Machine Learning, Heidelberg, Springer. Chapter 8 on graphical models openly available:
<http://research.microsoft.com/en-us/um/people/cmbishop/prml/>




Murphy, K. P. 2012. Machine learning: a probabilistic perspective, MIT press. Chapter 26 (pp. 907) – Graphical model structure learning




Koller, D. & Friedman, N. 2009. Probabilistic graphical models: principles and techniques, MIT press.




Rubinstein, R. Y. & Kroese, D. P. 2013. The cross-entropy method: a unified approach to combinatorial optimization, Monte-Carlo simulation and machine learning, Springer




Cameron Davidson-Pilon 2015. Bayesian methods for hackers: probabilistic programming and Bayesian inference, Addison-Wesley Professional.



Rubinstein, R. Y. & Kroese, D. P. 2013. Simulation and the Monte-Carlo Method, Wiley




Stiller, A., Goodman, N. & Frank, M. C. Ad-hoc scalar implicature in adults and children. CogSci, 2011.



<https://goo.gl/6a7rOC>

Chapter 8 Graphical Models is as sample chapter fully downloadable for free

Bishop, C. M. 2006. Pattern Recognition and Machine Learning, Heidelberg, Springer.



<http://bayes.cs.ucla.edu/BOOK-2K/>

Pearl, J. 2009. Causality: Models, Reasoning, and Inference (2nd Edition), Cambridge, Cambridge University Press.

- Medicine is an extremely complex application domain – dealing most of the time with uncertainties -> **probable information!**
- Key: Structure learning and prediction in large-scale biomedical networks with probabilistic graphical models
- Causality and Probabilistic Inference
- Uncertainties are present at all levels in health related systems
- Data sets from which ML learns are noisy, mislabeled, atypical, etc. etc.
- Even with data of high quality, gauging and combining a multitude of data sources and constraints in usually imperfect models of the world requires us to represent and process **uncertain knowledge** in order to make **viable decisions in context and within reasonable time!**
- In the increasingly complicated settings of modern science, model structure or causal relationships may not be known a-priori [1].
- Approximating probabilistic inference in Bayesian belief networks is NP-hard [2] -> here we need the “human-in-the-loop” [3]

[1] Sun, X., Janzing, D. & Schölkopf, B. Causal Inference by Choosing Graphs with Most Plausible Markov Kernels. ISAIM, 2006.

[2] Dagum, P. & Luby, M. 1993. Approximating probabilistic inference in Bayesian belief networks is NP-hard. Artificial intelligence, 60, (1), 141-153.

[3] Holzinger, A. 2016. Interactive Machine Learning for Health Informatics: When do we need the human-in-the-loop? Springer Brain Informatics (BRIN), 3, 1-13, doi:10.1007/s40708-016-0042-6.

- Reinforcement Learning is the **oldest approach**, with the longest history and can provide insight into understanding human learning [1]
- RL is the “**AI problem in the microcosm**” [2]
- Future opportunities are in Multi-Agent RL (MARL), Multi-Task Learning (MTL), Generalization and **Transfer-Learning** [3], [4].

[1] Turing, A. M. 1950. Computing machinery and intelligence. *Mind*, 59, (236), 433-460.

[2] Littman, M. L. 2015. Reinforcement learning improves behaviour from evaluative feedback. *Nature*, 521, (7553), 445-451, doi:10.1038/nature14540.

[3] Taylor, M. E. & Stone, P. 2009. Transfer learning for reinforcement learning domains: A survey. *The Journal of Machine Learning Research*, 10, 1633-1685.

[4] Pan, S. J. & Yang, Q. A. 2010. A Survey on Transfer Learning. *IEEE Transactions on Knowledge and Data Engineering*, 22, (10), 1345-1359, doi:10.1109/tkde.2009.191.

- **I) Supervised learning (classification)**
 - $y = f(x)$
 - Given x, y pairs; find a f that map a new x to a proper y
 - Regression, logistic regression, classification
 - Expert provides examples e.g. classification of clinical images
 - Disadvantage: Supervision can be expensive
- **II) Unsupervised learning (clustering)**
 - $f(x)$
 - Given x (features only), find f that gives you a description of x
 - Find similar points in high-dim X
 - E.g. clustering of medical images based on their content
 - Disadvantage: Not necessarily task relevant
- **III) Reinforcement learning**
 - $y = f(x)$
 - more general than supervised/unsupervised learning
 - learn from interaction to achieve a goal
 - Learning by direct interaction with environment (automatic ML)
 - Disadvantage: broad difficult approach, problem with high-dim data

JOURNAL OF THE AMERICAN STATISTICAL ASSOCIATION

Number 247

SEPTEMBER 1949

Volume 44

THE MONTE CARLO METHOD

NICHOLAS METROPOLIS AND S. ULMAM

Los Alamos Laboratory

We shall present here the motivation and a general description of a method dealing with a class of problems in mathematical physics. The method is, essentially, a statistical approach to the study of differential equations, or more generally, of integro-differential equations that occur in various branches of the natural sciences.

ALREADY in the nineteenth century a sharp distinction began to appear between two different mathematical methods of treating physical phenomena. Problems involving only a few particles were studied in classical mechanics, through the study of systems of ordinary differential equations. For the description of systems with very many particles, an entirely different technique was used, namely, the method of statistical mechanics. In this latter approach, one does not concentrate on the individual particles but studies the properties of *sets of particles*. In pure mathematics an intensive study of the properties of sets of points was the subject of a new field. This is the so-called theory of sets, the basic theory of integration, and the twentieth century development of the theory of probabilities prepared the formal apparatus for the use of such models in theoretical physics, i.e., description of properties of aggregates of points rather than of individual points and



Image Source:

<http://www.manhattanprojectvoices.org/oral-histories/nicholas-metropolis-interview>

THE JOURNAL OF CHEMICAL PHYSICS

VOLUME 21, NUMBER 6

JUNE, 1953

Equation of State Calculations by Fast Computing Machines

NICHOLAS METROPOLIS, ARIANNA W. ROSENBLUTH, MARSHALL N. ROSENBLUTH, AND AUGUSTA H. TELLER,
Los Alamos Scientific Laboratory, Los Alamos, New Mexico

AND

EDWARD TELLER,* *Department of Physics, University of Chicago, Chicago, Illinois*

(Received March 6, 1953)

A general method, suitable for fast computing machines, for investigating such properties as equations of state for substances consisting of interacting individual molecules is described. The method consists of a modified Monte Carlo integration over configuration space. Results for the two-dimensional rigid-sphere system have been obtained on the Los Alamos MANIAC and are presented here. These results are compared to the free volume equation of state and to a four-term virial coefficient expansion.

I. INTRODUCTION

THE purpose of this paper is to describe a general method, suitable for fast electronic computing machines, of calculating the properties of any substance which may be considered as composed of interacting individual molecules. Classical statistics is assumed, only two-body forces are considered, and the potential field of a molecule is assumed spherically symmetric. These are the usual assumptions made in theories of liquids. Subject to the above assumptions, the method is not restricted to any range of temperature or density. This paper will also present results of a preliminary two-dimensional calculation for the rigid-sphere system. Work on the two-dimensional case with a Lennard-Jones potential is in progress and will be reported in a later paper. Also, the problem in three dimensions is being investigated.

II. THE GENERAL METHOD FOR AN ARBITRARY POTENTIAL BETWEEN THE PARTICLES

In order to reduce the problem to a feasible size for numerical work, we can, of course, consider only a finite number of particles. This number N may be as high as several hundred. Our system consists of a square† containing N particles. In order to minimize the surface effects we suppose the complete substance to be periodic, consisting of many such squares, each square containing N particles in the same configuration. Thus we define d_{AB} , the minimum distance between particles A and B , as the shortest distance between A and any of the particles B , of which there is one in each of the squares which comprise the complete substance. If we have a potential which falls off rapidly with distance, there will be at most one of the distances AB which can make a substantial contribution; hence we need consider only the minimum distance d_{AB} .

* Now at the Radiation Laboratory of the University of California, Livermore, California.

† We will use the two-dimensional nomenclature here since it is easier to visualize. The extension to three dimensions is obvious.

Metropolis, N., Rosenbluth, A. W., Rosenbluth, M. N., Teller, A. H. & Teller, E. 1953. Equation of State Calculations by Fast Computing Machines. *The Journal of Chemical Physics*, 21, (6), 1087-1092, doi:10.1063/1.1699114.

Hastings, W. K. 1970. Monte Carlo sampling methods using Markov chains and their applications. *Biometrika*, 57, (1), 97-109.

Biometrika (1970), 57, 1, p. 97

Printed in Great Britain

97

Monte Carlo sampling methods using Markov chains and their applications

By W. K. HASTINGS

University of Toronto

SUMMARY

A generalization of the sampling method introduced by Metropolis *et al.* (1953) is presented along with an exposition of the relevant theory, techniques of application and methods and difficulties of assessing the error in Monte Carlo estimates. Examples of the methods, including the generation of random orthogonal matrices and potential applications of the methods to numerical problems arising in statistics, are discussed.

1. INTRODUCTION


For numerical problems in a large number of dimensions, Monte Carlo methods are often more efficient than conventional numerical methods. However, implementation of the Monte Carlo methods requires sampling from high dimensional probability distributions and this may be very difficult and expensive in analysis and computer time. General methods for sampling from, or estimating expectations with respect to, such distributions are as follows.

(i) If possible, factorize the distribution into the product of one-dimensional conditional distributions from which samples may be obtained.


(ii) Use importance sampling, which may also be used for variance reduction. That is, in order to evaluate the integral

$$J = \int f(x)p(x)dx = E_p(f),$$

where $p(x)$ is a probability density function, instead of obtaining independent samples x_1, \dots, x_N from $p(x)$ and using the estimate $\hat{J}_N = \sum f(x_i)/N$, we instead obtain the sample from



(a)



(b)

Henao, R., Lu, J. T., Lucas, J. E., Ferranti, J. & Carin, L. 2016. Electronic health record analysis via deep poisson factor models. Journal of Machine Learning Research JMLR, 17, 1-32.

Graphical Model Learning

- Remember: GM are a marriage between probability theory and graph theory and provide a tool for dealing with our two grand challenges in the biomedical domain:

Uncertainty and complexity


- The learning task is two-fold:
 - 1) Learning unknown probabilities
 - 2) Learning unknown structures

Jordan, M. I. 1998. Learning in graphical models, Springer

- 1) Test if a distribution is decomposable with regard to a given graph.
 - This is the most direct approach. It is not bound to a graphical representation,
 - It can be carried out w.r.t. other representations of the set of subspaces to be used to compute the (candidate) decomposition of a given distribution.
- 2) Find a suitable graph by measuring the strength of dependences.
 - This is a heuristic, but often highly successful approach, which is based on the frequently valid assumption that in a conditional independence graph an attribute is more strongly dependent on adjacent attributes than on attributes that are not directly connected to them.
- 3) Find an independence map by conditional independence tests.
 - This approach exploits the theorems that connect conditional independence graphs and graphs that represent decompositions.
 - It has the advantage that a single conditional independence test, if it fails, can exclude several candidate graphs. Beware, because wrong test results can thus have severe consequences.


Borgelt, C., Steinbrecher, M. & Kruse, R. R. 2009. Graphical models: representations for learning, reasoning and data mining, John Wiley & Sons.

Who of you is
NON-Smoker ?




Normal airway

Who of you is
Smoker ?



Asthmatic airway




Asthmatic airway
during attack

Beasley, R. 1998. Worldwide variation in prevalence of symptoms of asthma, allergic rhinoconjunctivitis, and atopic eczema: ISAAC. *The Lancet*, 351, (9111), 1225-1232,
doi:[http://dx.doi.org/10.1016/S0140-6736\(97\)07302-9](http://dx.doi.org/10.1016/S0140-6736(97)07302-9).



Bayesian Network

Patient	J46	Tussis	Smoker
Florian	1	1	0
Tamas	0	0	0
Matthias	1	0	0
Benjamin	0	1	1
Dimitrios	0	1	0
...			
...			




Florian	0	?	?
---------	---	---	---

Florian	0	0.3	0.2
---------	---	-----	-----

Rows are independent during learning and inference!

- Asthma can be hereditary
- Friends may have similar smoking habits
- Augmenting graphical model with relations between the entities – Markov Logic



2.1 $\text{Asthma} \Rightarrow \text{Cough}$


3.5 $\text{Smokes} \Rightarrow \text{Cough}$

2.1 $\text{Asthma}(x) \Rightarrow \text{Cough}(x)$

3.5 $\text{Smokes}(x) \Rightarrow \text{Cough}(x)$

1.9 $\text{Smokes}(x) \wedge \text{Friends}(x,y) \Rightarrow \text{Smokes}(y)$

1.5 $\text{Asthma}(x) \wedge \text{Family}(x,y) \Rightarrow \text{Asthma}(y)$




Example for probabilistic rule learning, in which probabilistic rules are learned from probabilistic examples: The ProbFOIL+ Algorithm solves this problem by combining the principles of the rule learner FOIL with the probabilistic Prolog called ProbLog, see: De Raedt, L., Dries, A., Thon, I., Van Den Broeck, G. & Verbeke, M. 2015. Inducing probabilistic relational rules from probabilistic examples. International Joint Conference on Artificial Intelligence (IJCAI).



Protein network inference from multiple genomic data: a supervised approach

Y. Yamanishi^{1,*}, J.-P. Vert² and M. Kanehisa¹

¹Bioinformatics Center, Institute for Chemical Research, Kyoto University, Gokasho, Uji, Kyoto 611-0011, Japan and ²Computational Biology group, Ecole des Mines de Paris, 35 rue Saint-Honoré, 77305 Fontainebleau cedex, France



- K_{exp} (Expression)
- K_{ppi} (Protein interaction)
- K_{loc} (Localization)
- K_{phy} (Phylogenetic profile)
- $K_{\text{exp}} + K_{\text{ppi}} + K_{\text{loc}} + K_{\text{phy}}$
(Integration)



A statistical framework for genomic data fusion


Gert R. G. Lanckriet¹, Tijl De Bie³, Nello Cristianini⁴,
Michael I. Jordan² and William Stafford Noble^{5,*}

¹Department of Electrical Engineering and Computer Science, ²Division of Computer Science, Department of Statistics, University of California, Berkeley 94720, USA,

³Department of Electrical Engineering, ESAT-SCD, Katholieke Universiteit Leuven 3001, Belgium, ⁴Department of Statistics, University of California, Davis 95618, USA and


⁵Department of Genome Sciences, University of Washington, Seattle 98195, USA

Kernel	Data	Similarity measure
K_{SW}	protein sequences	Smith-Waterman
K_B	protein sequences	BLAST
K_{Pfam}	protein sequences	Pfam HMM
K_{FFT}	hydropathy profile	FFT
K_{LI}	protein interactions	linear kernel
K_D	protein interactions	diffusion kernel
K_E	gene expression	radial basis kernel
K_{RND}	random numbers	linear kernel



Lanckriet, G. R., De Bie, T., Cristianini, N., Jordan, M. I. & Noble, W. S. 2004. A statistical framework for genomic data fusion. Bioinformatics, 20, (16), 2626-2635.

Graph




Model




06 Graphical Models and Decision Making

Data

$$\mathcal{D} \equiv \{X_1^{(i)}, X_2^{(i)}, \dots, X_m^{(i)}\}_{i=1}^N$$



Murphy, K. P. 2012. Machine learning: a probabilistic perspective, Cambridge (MA), MIT press.



π ... multinomial parameter vector, Stationary distribution of Markov chain





Image credit to Anna Goldenberg, Toronto

Goal: Learn an **optimal policy** for selecting best actions within a given **context**



For $t = 1, \dots, T$

- 1) The world produces a “context” $x_t \in X$
- 2) The learner selects an action

$$a_t \in \{1, \dots, K\}$$

- 3) The world reacts with a reward $r_t(a_t) \in [0,1]$


- Key Idea: Conditional independence assumptions are very useful – however: Naïve Bayes is extreme!
- X is *conditionally independent* of Y , given Z , if the $P(X)$ governing X is independent of value Y , given value of Z :

$$(\forall i, j, k) P(X = x_i | Y = y_j, Z = z_k) = P(X = x_i | Z = z_k)$$

can be abbr. with $P(X|Y, Z) = P(X|Z)$



- Graphical models express sets of conditional independence assumptions via graph structure
- The graph structure plus associated parameters define joint probability distribution over the set of variables

- Medicine is an extremely complex application domain – dealing most of the time with uncertainties -> **probable information!**
- When we have big data but little knowledge automatic ML can help to gain insight:
- **Structure learning and prediction in large-scale biomedical networks with probabilistic graphical models**
- If we have little data and deal with NP-hard problems we still need the human-in-the-loop




Undirected: Markov random fields, useful e.g. for computer vision (Details: Murphy 19)

$$P(\mathbf{X}) = \frac{1}{Z} \exp \left(\sum_{ij} W_{ij} x_i x_j + \sum_i x_i b_i \right)$$




Directed: Bayes Nets, useful for designing models (Details: Murphy 10)

$$p(\mathbf{x}) = \prod_{k=1}^K p(x_k | \text{pa}_k)$$



Factored: useful for inference/learning

$$p(\mathbf{x}) = \prod_s f_s(\mathbf{x}_s)$$




Baldi, P. & Pollastri, G. 2003. The principled design of large-scale recursive neural network architectures--dag-rnns and the protein structure prediction problem. *The Journal of Machine Learning Research*, 4, 575-602.

- Hypothesis: most biological functions involve the interactions between many proteins, and the complexity of living systems arises as a result of such interactions.
- In this context, the problem of inferring a global protein network for a given organism,
- - using all (genomic) data of the organism,
- is one of the main challenges in computational biology

Yamanishi, Y., Vert, J.-P. & Kanehisa, M. 2004. Protein network inference from multiple genomic data: a supervised approach. Bioinformatics, 20, (suppl 1), i363-i370.

Borgwardt, K. M., Ong, C. S., Schönauer, S., Vishwanathan, S., Smola, A. J. & Kriegel, H.-P. 2005. Protein function prediction via graph kernels. Bioinformatics, 21, (suppl 1), i47-i56.



- Important for health informatics: Discovering relationships between biological components
- Unsolved problem in computer science:
- Can the graph isomorphism problem be solved in polynomial time?
 - So far, no polynomial time algorithm is known.
 - It is also not known if it is NP-complete
 - We know that subgraph-isomorphism is NP-complete